

HELLO DEVDAY

Justin Marney

HELLO DEVDAY

Viget Labs

Justin Marney

HELLO DEVDAY

GMU BCS '06

Viget Labs

Justin Marney

HELLO DEVDAY

Ruby '07

GMU BCS '06

Viget Labs

Justin Marney

Viget '08

Ruby '07

GMU BCS '06

Viget Labs





DISTRIBUTING YOUR DATA

WHY?

**DISTRIBUTING
YOUR DATA**

**web applications are judged
by their level of availability**

WHY?

**DISTRIBUTING
YOUR DATA**

**ability to continue operating
during failure scenarios**

web applications are judged
by their level of availability

WHY?

**DISTRIBUTING
YOUR DATA**

**ability to manage availability
during failure scenarios**

ability to continue operating
during failure scenarios

web applications are judged
by their level of availability

WHY?

increase throughput

ability to **manage availability**
during node failure

ability to continue operating
during **failure scenarios**

web applications are judged
by their level of availability

WHY?

increase durability

increase throughput

ability to **manage availability**
during node failure

ability to continue operating
during **failure scenarios**

web applications are judged
by their level of availability

increase scalability

increase durability

increase throughput

ability to **manage availability**
during node failure

ability to continue operating
during **failure scenarios**

SCALABILITY

"I can add twice as much **X** and get twice as much **Y**."

X = processor, RAM, disks, servers, bandwidth

Y = throughput, storage space, uptime

SCALABILITY

scalability is a **ratio**.

2:2 = linear scalability ratio

scalability ratio allows you to predict how much it will cost you to grow.

SCALABILITY

UP/DOWN/VERTICAL/
HORIZONTAL/L/R/L/R/**A**/
B/START



SCALABILITY

UP

grow your infrastructure

multiple data centers

higher bandwidth

faster machines

SCALABILITY

DOWN

shrink your infrastructure

mobile

set-top

laptop

SCALABILITY

VERTICAL

add to a single node

CPU

RAM

RAID

SCALABILITY

HORIZONTAL

add more nodes

distribute the load

commodity cost

limited only by capital

@gary_hustwit: Dear
**Twitter: when a World Cup
match is at the 90th
minute, you might want to
turn on a few more servers.**

ASYNCHRONOUS

A distributed transaction is bound by availability of all nodes.

ASYNCHRONOUS

A distributed transaction is bound by availability of all nodes.

$$(.99^1) = .99$$

$$(.99^2) = .98$$

$$(.99^3) = .97$$

ASYNCHRONOUS

Asynchronous systems operate without the concept of **global state.**

The concurrency model more accurately reflects the real world.

ASYNCHRONOUS

Asynchronous systems operate without the concept of **global state.**

The concurrency model more accurately reflects the real world.

What about my **ACID!?**

ACID

Atomic

Series of database operations either all occur, or nothing occurs.

Consistent

Transaction does not violate any integrity constraints during execution.

Isolated

Cannot access data that is modified during an incomplete transaction.

Durable

Transactions that have committed will survive permanently.

ACID

Defines a set of characteristics that aim to **ensure consistency.**

What happens when we realize that in order scale we need to distribute our data and handle asynchronous operations?

ACID

Without global state, no Atomicity.

Without a linear timeline, no transactions and no Isolation.

The canonical location of data might not exist, therefore no D.

Without A, I, or D, Consistency in terms of entity integrity is no longer guaranteed.

CAP Theorem

Eric Brewer @ 2000 Principles of Distributed Computing (PODC).

Seth Gilbert and Nancy Lynch published a formal proof in 2002.

CAP Acronym

Consistency: Multiple values for the same piece of data are not allowed.

Availability: If a non-failing node can be reached the system functions.

Partition-Tolerance: Regardless of packet loss, if a non-failing node is reached the system functions.

CAP Theorem

Consistency, Availability, Partition-Tolerance: Choose One...

CAP Theorem

Single node systems bound by CAP.

100% Partition-tolerant

100% Consistent

No Availability Guarantee

CAP Theorem

Multi-node systems bound by **CAP**.

CA : DT, 2PC, **ACID**

CP : Quorum, distributed databases

AP : Dynamo, no **ACID**

CAP Theorem

CAP doesn't say **AP** systems are the solution to your problem.

Not an **absolute decision**.

Most systems are a hybrid of **CA**, **CP**, & **AP**.

CAP Theorem

Understand the trade-offs and use that understanding to build a system that fails predictably.

Enables you to build a system that degrades gracefully during a failure.

BASE

Dan Pritchett

BASE: An ACID Alternative

**Associate for Computing Machinery
Queue, 2008**

BASE

BASE: An ACID Alternative

Basically Available

Soft State

Eventually Consistent

BASE

BASE: An ACID Alternative

Basically Available

Soft State

Eventually Consistent

Eventually Consistent

Rename to **Managed Consistency**.

Does not mean **probable** or **hopeful**
or **indefinite time in the future**.

Describes what happens during a
failure.

Eventually Consistent

During certain scenarios a decision must be made to either return inconsistent data or deny a request.

EC allows you control the level of consistency vs. availability in your application.

Eventually Consistent

In order to achieve availability in an asynchronous system, accept that failures are going to happen.

Understand failure points and know what you are willing to give up in order to achieve availability.

**How can we model the operations
we perform on our data to be
asynchronous & EC?**

Model system as a network of independent components.

Partition components along functional boundaries.

Don't interact with your data as one big global state.

This doesn't meant **every part of
your system must operate this way!**

Use **ACID 2.0 to help identify and
architect components than can.**

ACID 2.0

Associative

Order of operations does not change the result.

Commutative

Operations can be aggregated in any order.

Idempotent

Operation can be applied multiple times without changing the result.

Distributed

Operations are distributed and processed asynchronously.

OPS BROS

Incremental scalability

Homogeneous node responsibilities

Heterogeneous node capabilities

LINKS

[Base: An ACID Alternative](#)

[Into the Clouds on New Acid](#)

[Brewer's CAP theorem](#)

[Embracing Concurrency At Scale](#)

[Amazon's Dynamo](#)

ME

<http://sorescode.com>

<http://github.com/gotascii>

<http://spkr8.com/s/1>

[@vigemarn](#)